

Untitled1

September 19, 2020

0.0.1 NLP : Project 2 - Help Twitter Combat Hate Speech Using NLP and Machine Learning

0.0.2 Importing All Library Packages

```
In [20]: import re
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import string
import nltk
import warnings
warnings.filterwarnings("ignore", category=DeprecationWarning)
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import confusion_matrix, f1_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier

%matplotlib inline
```

0.0.3 Load the tweets file using read_csv function from Pandas package.

```
In [2]: df = pd.read_csv(r'E:\study\simpli\NLP_proj\TwitterHate.csv')
```

```
In [3]: df.head()
```

```
Out[3]:
```

	id	label	tweet
0	1	0	@user when a father is dysfunctional and is s...
1	2	0	@user @user thanks for #lyft credit i can't us...
2	3	0	bihday your majesty
3	4	0	#model i love u take with u all the time in ...
4	5	0	factsguide: society now #motivation

```
In [4]: df.shape, df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
```

```
Data columns (total 3 columns):
id      31962 non-null int64
label   31962 non-null int64
tweet   31962 non-null object
dtypes: int64(2), object(1)
memory usage: 749.2+ KB
```

```
Out[4]: ((31962, 3), None)
```

```
In [5]: df['label'].value_counts()
```


```
Out[5]: 0      29720
        1       2242
        Name: label, dtype: int64
```

```
In [7]: import seaborn as sns
import re
import matplotlib.pyplot as plt
% matplotlib inline


import seaborn as sns
from plotly.offline import download_plotlyjs, init_notebook_mode, plot, iplot
```

0.0.4 Exploratory Data Analysis about the tweets

```
In [9]: #Exploratory Data Analysis
df['length'] = df['tweet'].apply(len)
fig1 = sns.barplot('label', 'length', data = df, palette='PRGn')
plt.title('Average Word Length vs label')
plot = fig1.get_figure()
plot.savefig('Barplot.png')
```



```
In [10]: #bar graph to count positive negative label
fig2 = sns.countplot(x= 'label',data = df)
plt.title('Label Counts')
plot = fig2.get_figure()
plot.savefig('Count Plot.png')
```



```
In [16]: df['length'].head()
```

```
Out[16]: 0    102
         1    122
         2     21
         3     86
         4     39
         Name: length, dtype: int64
```


```
In [17]: def vectorization(table):
         #CountVectorizer will convert a collection of text documents to a matrix of tokens
         #Produces a sparse representation of the counts
         #Initialize
         vector = CountVectorizer()
         #We fit and transform the vector created
         frequency_matrix = vector.fit_transform(table.tweet)
         #Sum all the frequencies for each word
         sum_frequencies = np.sum(frequency_matrix, axis=0)
         #Now we use squeeze to remove single-dimensional entries from the shape of an array
         #the sum of frequencies.
         frequency = np.squeeze(np.asarray(sum_frequencies))
         #Now we get into a dataframe all the frequencies and the words that they correspond to
         frequency_df = pd.DataFrame([frequency], columns=vector.get_feature_names()).transpose()
         return frequency_df
```

```

In [18]: def graph(word_frequency, sent):
    labels = word_frequency[0][1:51].index
    title = "Word Frequency for %s" %sent
    #Plot the figures
    plt.figure(figsize=(10,5))
    plt.bar(np.arange(50), word_frequency[0][1:51], width = 0.8, color = sns.color_palette("magma", 50),
            edgecolor = "black", capsize=8, linewidth=1);
    plt.xticks(np.arange(50), labels, rotation=90, size=14);
    plt.xlabel("50 more frequent words", size=14);
    plt.ylabel("Frequency", size=14);
    #plt.title('Word Frequency for %s', size=18) %sent;
    plt.title(title, size=18)
    plt.grid(False);
    plt.gca().spines["top"].set_visible(False);
    plt.gca().spines["right"].set_visible(False);
    plt.show()

In [21]: word_frequency = vectorization(df).sort_values(0, ascending = False)
    graph(word_frequency, 'all')

```



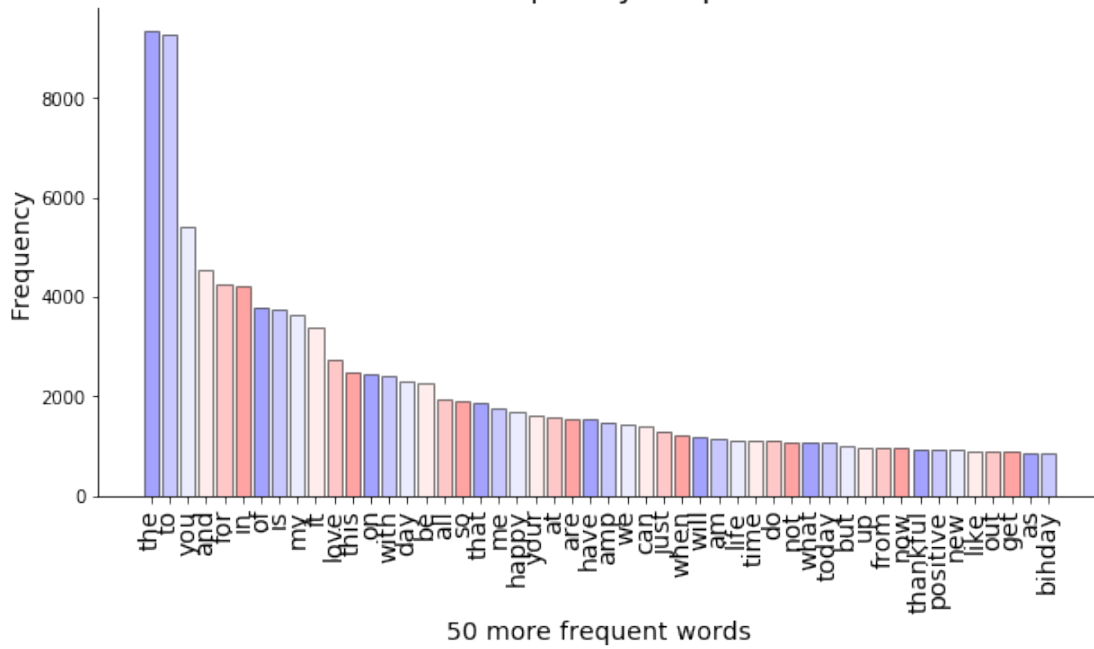
```

In [22]: word_frequency_pos = vectorization(df[df['label'] == 0]).sort_values(0, ascending = False)
    word_frequency_neg = vectorization(df[df['label'] == 1]).sort_values(0, ascending = False)


    graph(word_frequency_pos, 'positive')
    graph(word_frequency_neg, 'negative')

```

Word Frequency for positive



Word Frequency for negative



```
In [23]: def regression_graph(table):
          table = table[1:]
```

```

#We set the style of seaborn
sns.set_style("whitegrid")
#Initialize the figure
plt.figure(figsize=(6,6))

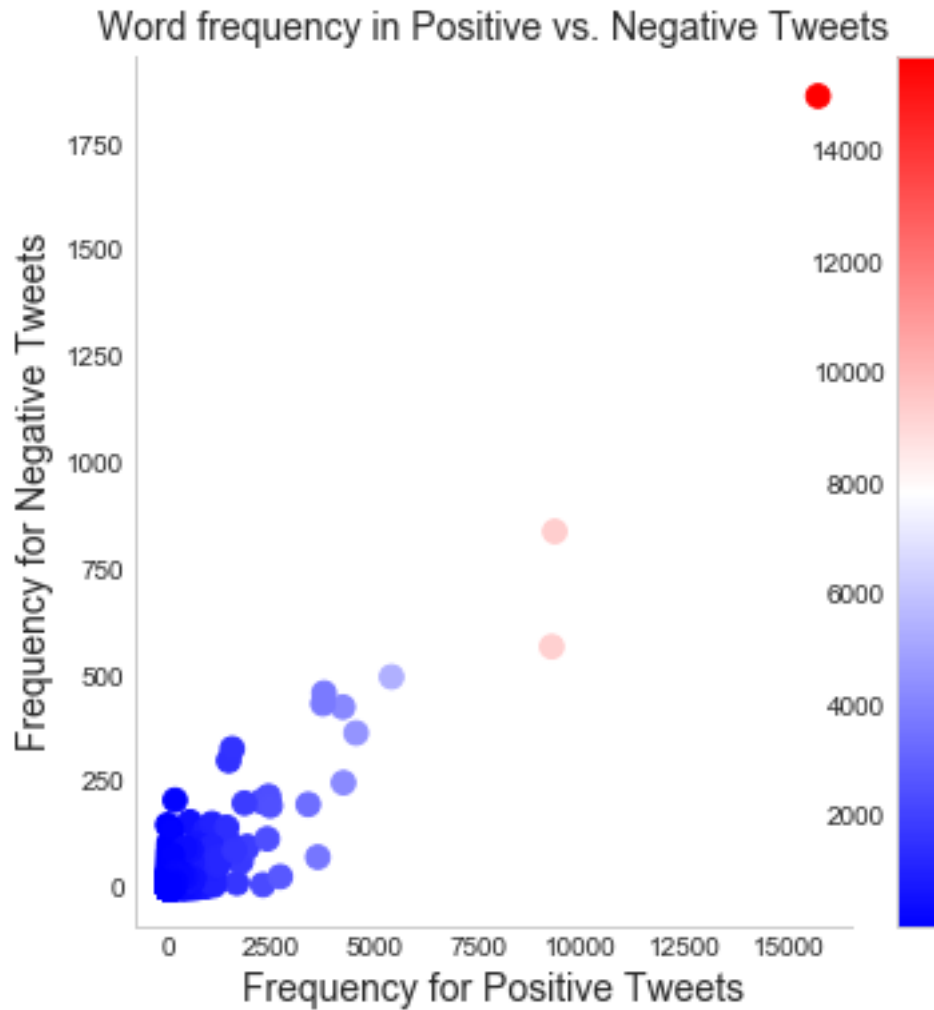
#we obtain the points from matplotlib scatter
points = plt.scatter(table["Positive"], table["Negative"], c=table["Positive"], s=
#graph the colorbar
plt.colorbar(points)
#we graph the regplot from seaborn
sns.regplot(x="Positive", y="Negative", fit_reg=False, scatter=False, color=".1", c
plt.xlabel("Frequency for Positive Tweets", size=14)
plt.ylabel("Frequency for Negative Tweets", size=14)
plt.title("Word frequency in Positive vs. Negative Tweets", size=14)
plt.grid(False)
sns.despine()

```

```

In [25]: table_regression = pd.concat([word_frequency_pos, word_frequency_neg], axis=1)
table_regression.columns = ["Positive", "Negative"]
regression_graph(table_regression)

```



0.0.5 Applying Preprocessing and Cleaning on the tweets

```
In [26]: def drop_features(features,data):
          data.drop(features,inplace=True,axis=1)
```

```
In [27]: import re
          ## example ##
          re.sub("(@[A-Za-z0-9]+)|([^0-9A-Za-z \t])", " ", "ouch...junior is angryð$ŸŘ#got7 #junior yugyo")
```

```
Out[27]: 'ouch junior is angry got7 junior yugyo '
```

```
In [28]: def process_tweet(tweet):
          return " ".join(re.sub("(@[A-Za-z0-9]+)|([^0-9A-Za-z \t])", " ",tweet.lower()).split())
```

```
In [29]: df['processed_tweets'] = df['tweet'].apply(process_tweet)
```

```
In [30]: df.head()
```



```

Out[30]:   id  label      tweet  length \
0    1    0  @user when a father is dysfunctional and is s...    102
1    2    0  @user @user thanks for #lyft credit i can't us...    122
2    3    0                                bihday your majesty     21
3    4    0  #model    i love u take with u all the time in ...     86
4    5    0                                factsguide: society now    #motivation    39

                                processed_tweets
0  when a father is dysfunctional and is so selfi...
1  thanks for lyft credit i can t use cause they ...
2                                bihday your majesty
3      model i love u take with u all the time in ur
4                                factsguide society now motivation

```

```
In [31]: drop_features(['id','tweet'],df)
```

```
In [32]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 31962 entries, 0 to 31961
Data columns (total 3 columns):
label                31962 non-null int64
length              31962 non-null int64
processed_tweets     31962 non-null object
dtypes: int64(2), object(1)
memory usage: 749.2+ KB

```

Train Test Split and Determining TF-IDF vectorizer values

```
In [33]: from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(df["processed_tweets"], df["label"],
```

```
In [34]: from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
count_vect = CountVectorizer(stop_words='english')
transformer = TfidfTransformer(norm='l2',sublinear_tf=True)
```

```
In [35]: x_train_counts = count_vect.fit_transform(x_train)
x_train_tfidf = transformer.fit_transform(x_train_counts)
```

```
In [36]: print(x_train_counts.shape)
print(x_train_tfidf.shape)
```

```

(25569, 33735)
(25569, 33735)

```

```
In [37]: x_test_counts = count_vect.transform(x_test)
x_test_tfidf = transformer.transform(x_test_counts)
```

```
In [38]: print(x_test_counts.shape)
        print(x_test_tfidf.shape)
```

```
(6393, 33735)
```

```
(6393, 33735)
```

0.0.6 Applying Model Implementation :

```
In [39]: from sklearn.linear_model import LogisticRegression
        from sklearn.naive_bayes import MultinomialNB
        from sklearn.ensemble import RandomForestClassifier
        from sklearn.linear_model import SGDClassifier
        from sklearn.metrics import f1_score
```

```
In [49]: #### Model building: Ordinary Logistic Regression
```

```
In [40]: modelLR = LogisticRegression(C=100).fit(x_train_tfidf,y_train)
```

```
In [50]: #### Model evaluation: Accuracy, recall, and f_1 score.
```

```
In [41]: predictionsLR = modelLR.predict(x_test_tfidf)
        sum(predictionsLR==1),len(y_test),f1_score(y_test,predictionsLR)
```

```
Out[41]: (334, 6393, 0.7063291139240505)
```

```
In [51]: #### Train again with the adjustment and evaluate :Regularization and Hyperparameter
```

```
In [42]: from sklearn.pipeline import Pipeline
        from sklearn.naive_bayes import MultinomialNB
        from sklearn.feature_extraction.text import CountVectorizer, TfidfTransformer
        from sklearn.model_selection import train_test_split, GridSearchCV
```

```
In [43]: text_clf = Pipeline([('vect', CountVectorizer()),
                              ('tfidf', TfidfTransformer()),
                              ('clf', MultinomialNB())])
        tuned_parameters = {
            'vect__ngram_range': [(1, 1), (1, 2), (2, 2)],
            'tfidf__use_idf': (True, False),
            'tfidf__norm': ('l1', 'l2'),
            'clf__alpha': [1, 1e-1, 1e-2]
        }
        from sklearn.metrics import classification_report
        clf = GridSearchCV(text_clf, tuned_parameters, cv=10)
        clf.fit(x_train, y_train)
```

```
Out[43]: GridSearchCV(cv=10, error_score='raise',
                      estimator=Pipeline(memory=None,
                      steps=[('vect', CountVectorizer(analyzer='word', binary=False, decode_error='str',
                      dtype=<class 'numpy.int64'>, encoding='utf-8', input='content',
```

```

lowercase=True, max_df=1.0, max_features=None, min_df=1,
ngram_range=(1, 1), preprocessor=None, stop_words=None,
strip...inear_tf=False, use_idf=True)), ('clf', MultinomialNB(alpha=1.0, clas
fit_params=None, iid=True, n_jobs=1,
param_grid={'vect_ngram_range': [(1, 1), (1, 2), (2, 2)], 'tfidf__use_idf': (
pre_dispatch='2*n_jobs', refit=True, return_train_score='warn',
scoring=None, verbose=0)

```

```
In [45]: print(classification_report(y_test, clf.predict(x_test), digits=4))
```

	precision	recall	f1-score	support
0	0.9666	0.9939	0.9801	5937
1	0.8750	0.5526	0.6774	456
avg / total	0.9601	0.9625	0.9585	6393

```
In [52]: ### we have received the best parameters and stats .....
```